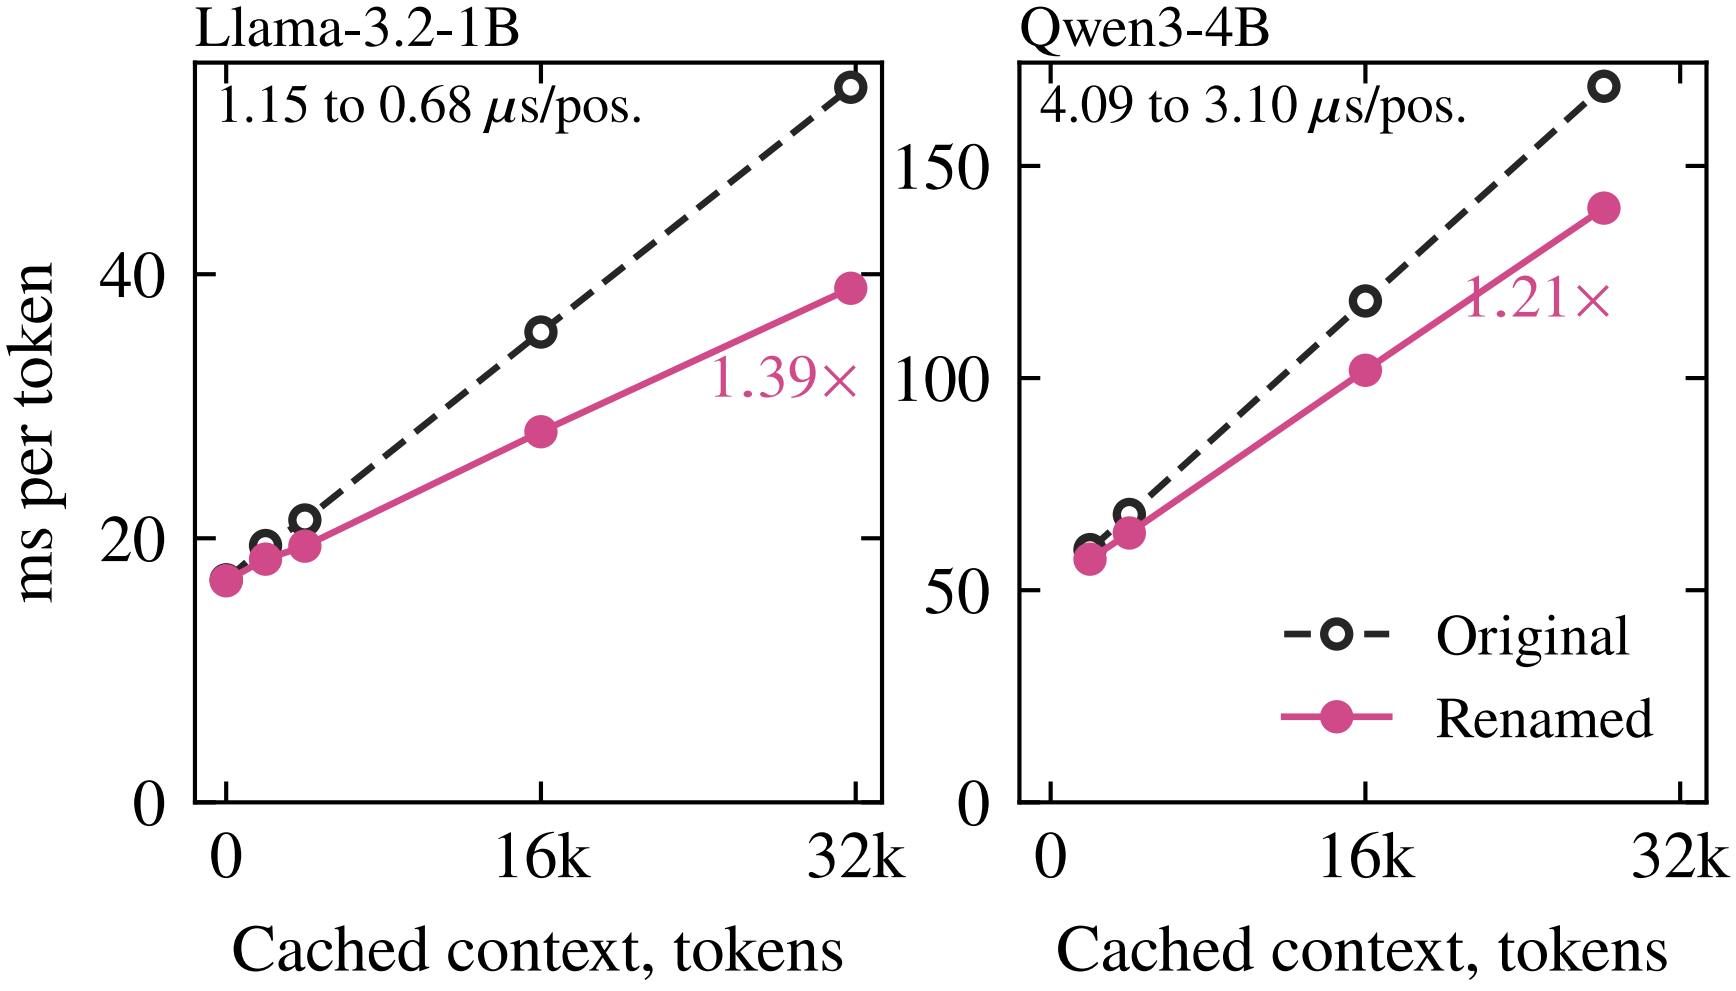

Llama-3.2-1B
1.15 to 0.68 μs/pos.
1.39×
Qwen3-4B
4.09 to 3.10 μs/pos.
1.21×
ms per token
0
20
40
0
50
100
150
0
16k
32k
Cached context, tokens
Original
Renamed